

Name: Xinrun Zhang

ML Homework 2 Report

Instructor: Prof. Kaliappan Gopalan


Time: 21/02/2019

1. Files

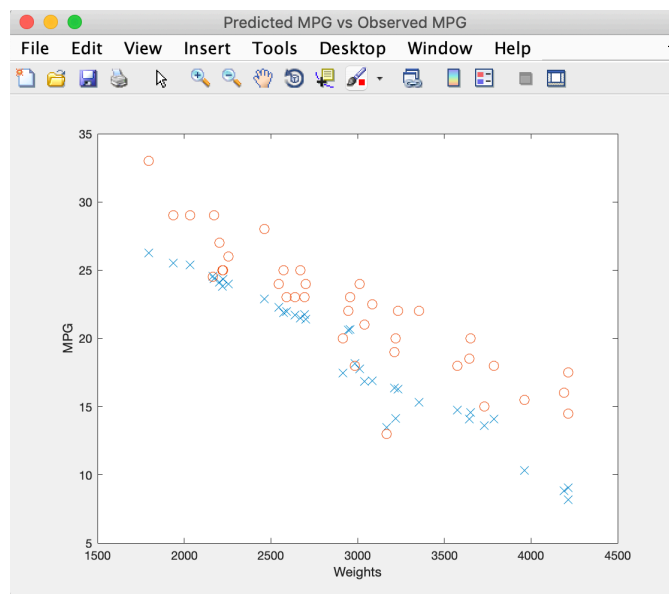
- a) `main.m` – Includes the main program of the project.
- b) `computeCost.m` – This function is used to compute the cost between hypothesis and input y .
- c) `normalizing.m` – This function is used in univariate linear regression to normalize the input x .
- d) `gradientDescent.m` – This function is used in univariate linear regression to execute the gradient descent to find proper parameters.

2. Multivariate Linear Regression with Gradient Descent


1. Scatter plot of y vs. x , where x is still the weight (as in HW 1)



2. Plot of predicted MPG against observed MPG for the validation data



3. Plot of J vs. Iteration index for the training data



4. Minimum J and the hypothesis parameters from gradient descent code

```
Theta found by gradient descent:
19.024179
-1.095249
-1.274129
-0.275286
-2.802676

With theta = [19.024179 ; -1.095249; -1.274129; -0.275286; -2.802676]
Cost computed = 2.925580
```

5. Mean squared error err_cv


```
The mean square error err_cv is: 17.800543
```

3. Multivariate Linear Regression with Normal Equation

1. Hypothesis parameters (theta) for the closed form

```
Theta found by normal equation:  
35.486442  
-0.296053  
-0.010934  
0.000843  
-0.003744  
  
With theta = [35.486442 ; -0.296053; -0.010934; 0.000843; -0.003744]  
Cost computed = 2.876415
```

2. Plot of predicted MPG against observed MPG for the validation data



3. I didn't apply feature scaling (normalization) because when I searched this question on Google I found:

Well, in the second article there is a sentence:

Note that before conducting linear regression, you should normalize the data. One way is $\frac{x_i - \text{mean}(x)}{\text{Range}(x)}$, and some use $sd(x)$ as the denominator. Both work.

But is not said that it applies specifically to normal equations. And in gradient descent section there is nothing said about normalization. So I suppose it was a small mistake to include that sentence in Normal Equation section instead of Gradient Descent.

Anyway Andrew Ng is pretty authoritative on machine learning topic, so you can rely on his words:

Using this formula does not require any feature scaling, and you will get an exact solution in one calculation: there is no 'loop until convergence' like in gradient descent.

The result of normal equation is not as good as I expected, there are maybe some method to get a better result.

4. Optional Exercise


1. Minimum J and the hypothesis parameters from gradient descent code

Optional exercise Theta found by gradient descent:


```
19.024179  
-1.166674  
-1.426639  
-0.665490  
-2.627616  
-0.597876
```

```
With theta = [19.024179 ; -1.166674; -1.426639; -0.665490; -2.627616; -0.597876]  
Cost computed = 2.886151
```

2. Plot of predicted MPG against observed MPG for the validation data



3. Plot of J vs. Iteration index for the training data



4. The result of 5 features LR is pretty closed to the situation that with 4 features LR. I believe that we can improve the result by adding more data to the training set instead of adding new features.

5. Codes

a) main.m

```
%% Machine Learning Homework 2

% Author: Xinrun Zhang

% Time: 02/20/2019 14:39

% =====

%% Initialization

clear ; close all; clc

% Import the data;

fprintf('Initializing...\n')

fprintf('Reading the data...\n');

A = xlsread('AutoData_HW2.xlsx');

% Extract the data, generate training data and validation data;
```

```

x = A(1:160,2:5);

y = A(1:160,1);

x_val = A(161:200,2:5);

y_val = A(161:200,1);


% Plot the data;

fprintf('Visualizing the data...\n\n')

figure('Name','Raw Data','NumberTitle','off');

plot(x(:,4),y,'x','MarkerSize',8);

ylabel('MPG');

xlabel('Weights');


% Initialize the theta vector;

theta = zeros(5, 1);


% Initialize the gradient descent parameters

iteration = 1000;

alpha = 0.01;

% =====

%% Data processing

fprintf('Data processing...\n')

% Call the normalizing function;

x_nor = normalizing(x);

x_val_nor = normalizing(x_val);

% Add a column of ones to xnor;

X = [ones(160, 1), x_nor(:,1:4)];

X_val = [ones(40, 1), x_val_nor(:,1:4)];

% =====

```

```

%% Multivariate linear regression

J = computeCost(X, y, theta); %compute the cost

fprintf('\nWith theta = [0 ; 0]\nCost computed = %f\n', J);

% Running gradient descent

theta = gradientDescent(X, y, theta, iteration, alpha);

% Print the output, including new theta and J;

fprintf('\nTheta found by gradient descent:\n');

fprintf('%f\n', theta);

J = computeCost(X, y, theta);

fprintf('\nWith theta = [%f ; %f; %f; %f; %f]\nCost computed = %f\n',
theta(1),theta(2),theta(3),theta(4),theta(5), J);

% =====

%% Validation

% Generate predict data

y_prd = X_val*theta;

% Calculate the mean squared error err_cv

err_cv = immse(y_prd,y_val);

fprintf('The mean sugare error err_cv is: %f\n',err_cv);

% Plot of predicted MPG against observed MPG for the validation data

figure('Name','Predicted MPG vs Observed MPG','NumberTitle','off');

plot(x_val(:,4),y_prd,'x','MarkerSize',8);

ylabel('MPG');

xlabel('Weights');

hold on;

plot(x_val(:,4),y_val,'o','MarkerSize',8);

% =====

%% Normal equation

```

```

% Data_processing

X_ne = [ones(160,1), x(:,1:4)];

X_ne_val = [ones(40,1), x_val(:,1:4)];

% Normal equation

theta_ne = pinv(X_ne)* y;

%theta_ne = (X'*X)^(-1)* X' * y;

y_prd_ne = X_ne_val*theta_ne;


% Print the output, including new theta and J;

fprintf('\nTheta found by normal equation:\n');

fprintf('%f\n', theta_ne);


J = computeCost(X_ne, y, theta_ne);

fprintf('\nWith theta = [%f ; %f; %f; %f; %f]\nCost computed = %f\n',
theta_ne(1),theta_ne(2),theta_ne(3),theta_ne(4),theta_ne(5), J);


% Plot of predicted MPG against observed MPG for the validation data

figure('Name','Predicted MPG vs Observed MPG 2','NumberTitle','off');

plot(x_val(:,4),y_prd_ne,'x','MarkerSize',8);

ylabel('MPG');

xlabel('Weights');

hold on;

plot(x_val(:,4),y_val,'o','MarkerSize',8);

% =====

%% Optional exercise

% Extract the data, generate training data and validation data;

x_oe = A(1:160,2:6);

y_oe = A(1:160,1);

x_oe_val = A(161:200,2:6);

y_oe_val = A(161:200,1);


% Initialize the theta vector;

```



```

theta_oe = zeros(6, 1);

% Initialize the gradient descent parameters

iteration = 1000;

alpha = 0.01;

% =====

% Data processing

x_oe_nor = normalizing(x_oe);

X_oe = [ones(160, 1), x_oe_nor(:,1:5)];

x_oe_val_nor = normalizing(x_oe_val);

X_oe_val = [ones(40, 1), x_oe_val_nor(:,1:5)];

% Gradient Descent

theta_oe = gradientDescent(X_oe, y, theta_oe, iteration, alpha);

% Print the output, including new theta and J;

fprintf('\nOptional exercise Theta found by gradient descent:\n');

fprintf('%f\n', theta_oe);

J = computeCost(X_oe, y_oe, theta_oe);

fprintf('\nWith theta = [%f ; %f; %f; %f; %f; %f]\nCost computed = %f\n',

theta_oe(1),theta_oe(2),theta_oe(3),theta_oe(4),theta_oe(5), theta_oe(6), J);

% =====

% Validation

% Generate predict data

y_prd = X_oe_val*theta_oe;

% Plot of predicted MPG against observed MPG for the validation data

figure('Name','Predicted MPG vs Observed MPG','NumberTitle','off');

plot(x_oe_val(:,4),y_prd,'x','MarkerSize',8);

ylabel('MPG');

```

```

xlabel('Weights');

hold on;

plot(x_oe_val(:,4),y_val,'o','MarkerSize',8);

% =====

```

b) computeCost.m

```

function J = computeCost(X, y, theta)

m = length(y);

J = 0;

hypothesis = X * theta;

J = (1/ (2*m)) * (hypothesis - y)' * (hypothesis - y);

end

```

c) gradientDescent.m

```

function [theta, J_history ] = gradientDescent(X, y, theta, iteration, alpha)

m = length(y);

J_history = zeros(iteration, 1);

for i = 1:iteration

    hypothesis = X*theta;

    theta = theta - (alpha / m) * (X' * (hypothesis - y));

    J_history(i) = computeCost(X, y, theta);

end

figure('Name','Multi LR: J vs Iteration Index','NumberTitle','off');

plot(J_history); % plot J vs iteration index

ylabel('J');

xlabel('Iteration');

end

```

d) normalizing.m

```
function x_nor = normalizing(x)

x_nor = x;

%delta = x_max - x_min;

x_mean = mean(x);

x_nor = (x - x_mean)./ (std(x));

end
```